



High-performance software - Easy gains with simple CUDA

Bernhard Raml

TU Wien
Department of Geodesy and Geoinformation
Research Area Remote Sensing
bernhard.raml@geo.tuwien.ac.at

Vienna | April 24, 2023

How to write high-performance software

Habitability comes before speed

- A habitable code base means automated tests!
- Design should focus on maintainability first.
- Measure using profiling tools to make informed decisions about *what to optimise*.
- Automated performance tests need a stable environment or cover only the most basic components.

Measure before you act

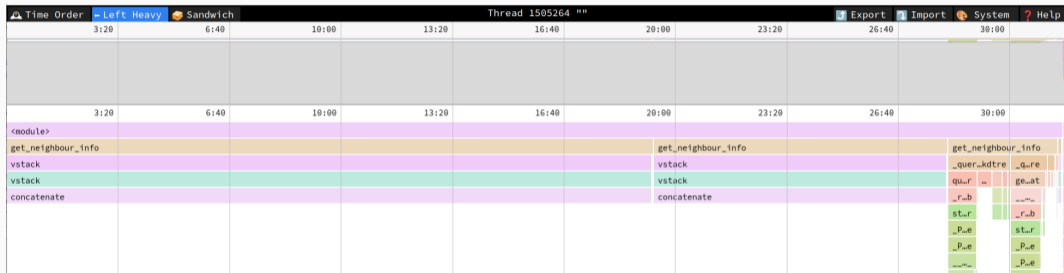


Figure 1: Repeated concatenation slow down due to mem copies

Measure before you act

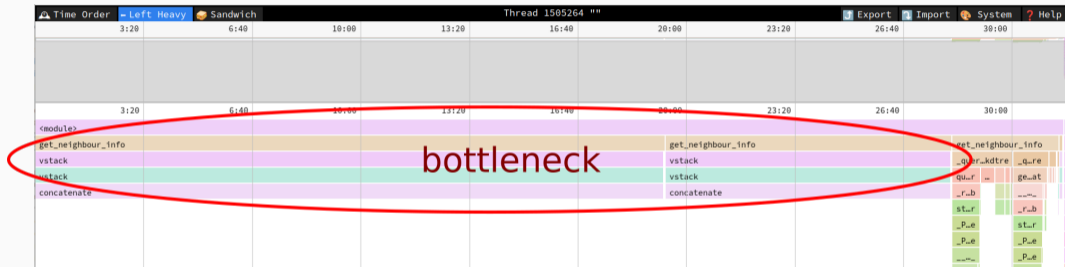


Figure 2: Repeated concatenation slow down due to mem copies

Streaming to increase throughput

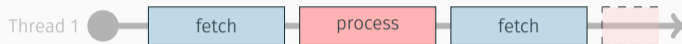


Figure 4: Serial process, waiting for data before processing




Figure 5: Stream next data-block while processing previous one

GPU vs CPU

The right tool for the right job

CPU

Complex control flow

Tree or graph search, sparse matrix operations

Serial processes

IO, compression

General purpose tasks

UI, web services, OS

GPU

Embarrassingly parallel tasks

Dense matrix multiplication


High memory throughput

Video processing, 3D rendering, high resolution remote sensing

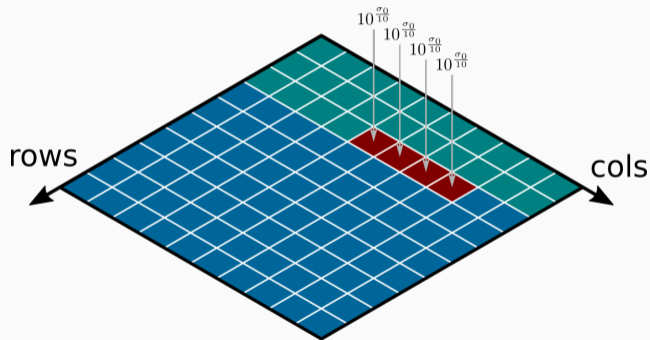
Specialized tasks

Ray-tracing, video codecs

The right tool for the right job - Examples



CPU - KD tree query



GPU - dB to linear

CPU

- Maximise instructions per cycle
- Low latency of single core
- Deeper cache hierarchy
- Complex instructions

GPU

- Maximise total throughput
- Streamed processes to hide higher latency
- Shallower cache hierarchy
- Simple instructions

The right tool for the right job



CPU

Image: Jorge Zapata - unsplash.com

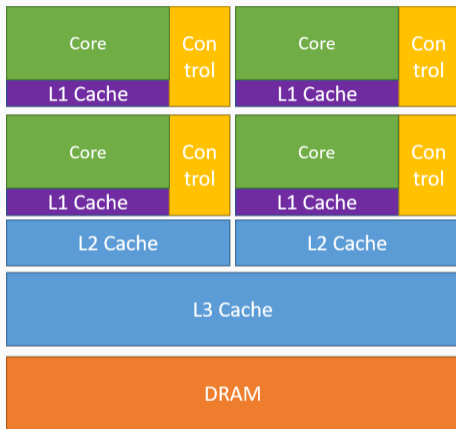


GPU


Image: pastaproductionline.com

GPU Architecture - Opening the magic box

Hardware Layout CPU vs. GPU



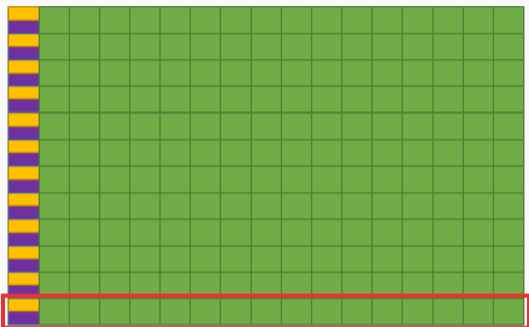
CPU



GPU

Image: CUDA Programming Guide

Memory latency - One cycle one second metaphor



Streaming Multiprocessors with shared memory



L2 Cache



DRAM

Register: 1s

Shared Memory: 2s - 4s

L1 Cache: 10s - 30s

L2 Cache: 1min

Device RAM: 2min - 10min

Host to device transfer: hours

Coding Session

Where to go from here?

- Of course, streaming from CPU-RAM to GPU-RAM (VRAM) improves throughput as well
- Avoid stalls from branching using clever distributing across *warps* or the step function trick
- Exploit specialised hardware accelerated *intrinsic* functions, e.g., add-mul
- Use different floating point representations like 16-bit halves
- Look into additional libraries within the CUDA ecosystem, e.g., cuBLAS, cuSolver...

Appendix

Davey Farley's YouTube Channel *Continuous Delivery* - Hardware cycles:

<https://www.youtube.com/watch?v=0reMVgn6kRo>

Wong, Henry, et al. "Demystifying GPU microarchitecture through microbenchmarking."

2010 IEEE International Symposium on Performance Analysis of Systems Software

CUDA Programming Guide:

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

Peng Wang's Fundamental Optimizations in CUDA Presentation:

https://developer.download.nvidia.com/GTC/PDF/1083_Wang.pdf

py-spy: <https://github.com/benfred/py-spy>

PyResample: <https://github.com/pytroll/pyresample>

Python Approval Tests:

<https://github.com/approvals/ApprovalTests.Python>

Special thanks to Raphael Quast for the Latex template